

# Automatic Cache Aware Roofline Model Building and Validation Using Topology Detection

Nicolas Denoyelle - Aleksandar Ilic - Brice Goglin - Leonel Sousa -  
Emmanuel Jeannot


Inria (France) - INESC-ID (Portugal) - Inria (France) - INESC-ID (Portugal) - Inria (France)

October 3, 2016

- 1 Introduction
- 2 The Cache Aware Roofline Model
- 3 Tool and Validation
- 4 Conclusion & Future Work

# Hardware Complexity is Growing

- ↗ Cores: up to 72 on Intel KNL.
- ↘ Memory per Core.
- ↗ Cache Sharing .
- ↗ Cache Hierarchy: several levels, side cache *etc.* . .



- 
- 1 Introduction
  - 2 The Cache Aware Roofline Model**
  - 3 Tool and Validation
  - 4 Conclusion & Future Work

# The Cache Aware Roofline Model (CARM)


The Cache Aware Roofline Model (CARM) simplifies this for you !

- What is the best you can get from the chip ?  
[Platform Model]
- Is it worth investing in code optimization ?  
[Application Model]

## Improvement over the Original Roofline Model (ORM).

### Assumptions:


- The machine is bound whether by the memory unit or by the compute unit.
- Both unit are not dependent on each other.




$$ddot = ddot + b[i] * c[i] \quad (1)$$

- memory:  $b[i], c[i] = 8 + 8\text{Bytes}$
- compute:  $+, * = 2\text{flops}$
- arithmetic intensity:  $2^{-3}\text{Flops/Bytes}$
- performance:  $\text{Flops/Time}$

# ORM versus CARM



Bytes transferred from DRAM to last level cache



Bytes issued from Cores.



- 1 Introduction
- 2 The Cache Aware Roofline Model
- 3 Tool and Validation
- 4 Conclusion & Future Work

# A Tool to Build and Validate the Model

## Build Model from Portable Blocks

- Hwloc gives the machine informations: cache structure and attributes, cores.
- The compiler gives architecture type and let us select the best instruction set.

## Micro-Benchmark Each Unit

- Bandwidth benchmark each cache level, and several stream types: Load, Store, 2Loads/Store.
- Floating Point peak benchmark: MUL, ADD, MAD, FMA.
- Micro Benchmarks of several arithmetic intensities for validation.


# A Tool for Non Experts

## Easy steps

- Build the model for the target platform.
- Collect CARM metrics from your application with our library (Hardware counters: PAPI).
- Plot this on a chart and get insights on your application.

## Quick Glance if the Model is Wrong

- Validation matches model ?
- Micro-Benchmarks matches theoretical throughput ?



- 
- 1 Introduction
  - 2 The Cache Aware Roofline Model
  - 3 Tool and Validation
  - 4 Conclusion & Future Work**

## Conclusion

We provide a Tool to build the CARM and model applications. It is portable and easy to use. It reaches near theoretical architecture throughput.

## Future Work

The CARM gives insight over Caches (i.e temporal locality). We will use it to model system heterogeneous memory, and give insights on data locality.

Thank you !

<https://github.com/NicolasDenoyelle/LARM-Locality-Aware-Roofline-Model->

# Automatic Cache Aware Roofline Model Building and Validation Using Topology Detection

Nicolas Denoyelle - Aleksandar Ilic - Brice Goglin - Leonel Sousa -  
Emmanuel Jeannot

Inria (France) - INESC-ID (Portugal) - Inria (France) - INESC-ID (Portugal) - Inria (France)

October 3, 2016